Hybrid Dense-Sparse Retrieval for High-Recall Information Retrieval

Abraham Itzhak Weinberg
AI-WEINBERG, AI Experts
aviw2010@gmail.com

Abstract

While dense retrieval methods have achieved state-of-the-art performance on ranking metrics, they often struggle with recall, particularly when queries require finding multiple relevant documents. We propose a simple yet effective hybrid approach that combines dense semantic embeddings with sparse lexical matching through weighted score fusion. Our method achieves up to 580% improvement in Recall@10 over dense-only retrieval on MS MARCO (13.9% $\rightarrow$ 80.8%), with strong performance maintained across multiple benchmarks. Through comprehensive experiments on three diverse datasets (MS MARCO, BEIR SciFact, BEIR NFCorpus) and comparisons with seven retrieval methods including ColBERT-style late interaction and learned sparse approaches, we demonstrate that simple hybrid retrieval is particularly effective for applications requiring high recall. We analyze dataset characteristics that determine when hybridization helps, provide parameter sensitivity analysis, and establish that a weighted combination ($\alpha = 0.5$) achieves optimal recall while maintaining acceptable ranking quality. Our findings suggest that hybrid approaches should be the default choice for recall-critical applications such as legal search, medical literature review, and comprehensive research tasks.

Keywords: Information Retrieval, Hybrid Retrieval, Dense Retrieval, Sparse Retrieval, Retrieval-Augmented Generation (RAG), BM25, Neural Retrieval, Semantic Search, Recall Optimization, Vector Search

1 Introduction

Information retrieval systems are increasingly adopting dense neural retrieval methods [5], which encode queries and documents into semantic embeddings and perform nearest neighbor search. These methods have demonstrated superior performance on ranking metrics like NDCG@10 and MRR compared to traditional sparse methods like BM25 (Best Matching 25) [15]. However, dense retrieval exhibits a critical weakness: **recall**.

Recall measures the fraction of relevant documents retrieved in the top-k results, making it crucial for applications where comprehensive coverage matters more than perfect ranking. Examples include:

- **Legal search**: Finding all precedents relevant to a case
- **Medical literature review**: Identifying all studies on a treatment
- **Patent search**: Discovering all prior art
- **Academic research**: Comprehensive literature surveys
- **Fact-checking and verification**: Gathering all evidence

While dense methods excel at semantic matching and understanding conceptual relationships, they can miss documents that contain exact query terms but lack overall semantic similarity. Conversely, sparse methods like BM25 excel at lexical matching and capturing exact terminology but fail to understand paraphrases or conceptual relationships. This complementarity suggests a hybrid approach combining both paradigms could achieve superior recall.

Hybrid Retrieval-Augmented Generation (Hybrid RAG) represents an architectural pattern that blends multiple retrieval approaches within a single pipeline [12]. Unlike pure vector-based semantic search or rule-driven keyword queries, Hybrid RAG capitalizes on the strengths of each approach—semantic understanding from dense retrieval and exact keyword matching from sparse retrieval—to deliver more accurate and resilient results [13].

Our Contributions:

1. We demonstrate that simple weighted fusion of dense and sparse retrieval achieves up to **580% improvement** (5.8×) in Recall@10 on MS MARCO over dense-only retrieval.

2. We provide comprehensive evaluation across three diverse datasets with varying characteristics, showing consistent improvements on passage-style queries but revealing dataset-specific behaviors.

3. We compare against seven retrieval methods including ColBERT-style late interaction [6] and learned sparse representations [3], establishing hybrid's competitive position.

4. We analyze dataset characteristics that determine when hybridization helps, identifying vocabulary diversity and query complexity as key factors.

5. We provide parameter sensitivity analysis across multiple α values, establishing $\alpha = 0.5$ as optimal for recall-focused scenarios.

6. We characterize the recall-NDCG trade-off quantitatively, enabling informed architectural decisions based on application requirements.

2 Related Work

2.1 Dense Retrieval

Dense retrieval methods encode text into continuous vector representations using neural networks. Dense Passage Retrieval (DPR) [5] pioneered bi-encoder architectures for passage retrieval, achieving strong results on open-domain Question Answering (QA). Subsequent work has explored improved training strategies [18], hard negative mining [19], and cross-encoder distillation [4]. While these methods excel at semantic matching, they struggle with exact term matching and vocabulary mismatch—critical for recall-oriented scenarios.

2.2 Sparse Retrieval

Classical sparse retrieval methods like BM25 [15] remain competitive baselines due to their efficiency and precision on exact matches. Recent work has explored learned sparse representations that use neural networks to learn term importance weights. SParse Lexical And Expansion (SPLADE) [3] expands documents with learned terms, achieving improvements over BM25 while maintaining sparse representations. Deep Contextualized Term weighting (DeepCT) [1] and High-Dimensional Contextualized Terms (HDCT) [2] learn contextualized term weights. These methods bridge sparse and dense approaches but add significant computational overhead compared to traditional BM25.

2.3 Late Interaction Models

Contextualized Late Interaction over BERT (ColBERT) [6] introduced late interaction, computing token-level similarities between queries and documents rather than single-vector representations. This approach achieves better recall than standard dense retrieval by capturing fine-grained matching signals. However, late interaction requires storing token embeddings for all documents, significantly increasing storage costs and retrieval latency compared to single-vector approaches.

2.4 Hybrid Retrieval Systems

Hybrid retrieval combines multiple retrieval paradigms to leverage complementary strengths. In the Retrieval-Augmented Generation (RAG) literature [13], hybrid approaches are distinguished from simpler single-method systems by their integration of both semantic (dense) and lexical (sparse) retrieval. Recent work by Luan et al. [9] combines sparse and dense representations during training. Wang et al. [17] use linear interpolation of scores for re-ranking. Ma et al. [10] study reproduction of hybrid systems in production environments.

Nigam [12] characterizes Hybrid RAG architectures along several dimensions: (1) *parallel retrieval*, where multiple indexes are queried simultaneously; (2) *cascading retrieval*, where broad semantic search is refined with keyword filters; and (3) *weighted scoring*, where different confidence scores are assigned to each retrieval source. Our work

focuses on the weighted scoring approach with parallel retrieval, as it provides the best balance between simplicity and effectiveness.

The broader RAG taxonomy [13] identifies four main categories: Simple RAG (single retrieval method), Hybrid RAG (combining retrieval techniques), Agentic RAG (multi-agent systems with specialized retrievers), and Graph RAG (leveraging graph structures). Our work falls squarely in the Hybrid RAG category but distinguishes itself through: (1) comprehensive evaluation on standard benchmarks, (2) quantitative analysis of the recall-ranking trade-off, and (3) characterization of dataset properties that determine hybrid effectiveness.

Unlike prior work that focuses on ranking improvements, we explicitly target recall-critical applications and provide systematic analysis of when and why hybrid approaches succeed or fail.

2.5 Retrieval-Augmented Generation

RAG systems [8] combine retrieval with generation to ground language model outputs in external knowledge. Most RAG systems prioritize ranking quality—measured by Normalized Discounted Cumulative Gain (NDCG) and Mean Reciprocal Rank (MRR)—over recall. However, recent work emphasizes the importance of comprehensive retrieval. The Factuality, Retrieval, And reasoning MEasurement Set (FRAMES dataset) [7] demonstrates that multi-hop reasoning requires high-recall retrieval to gather all necessary evidence. Our work provides empirical evidence that simple hybrid approaches can dramatically improve recall, making RAG systems more suitable for knowledge-intensive tasks.

3 Hybrid Retrieval: Definition and Architecture

3.1 Formal Definition

We define **Hybrid Retrieval** as a retrieval system that combines multiple retrieval paradigms through score fusion or ranking aggregation. Formally, given a query q and document collection $\mathcal{D} = \{d_1, \ldots, d_n\}$, a hybrid retriever computes:

$$s_{\text{hybrid}}(q, d_i) = f(s_1(q, d_i), s_2(q, d_i), \ldots, s_m(q, d_i)) \tag{1}$$

where $s_j(q, d_i)$ is the score from retrieval method $j \in \{1, \ldots, m\}$, and f is a fusion function. Common fusion strategies include weighted linear combination, Reciprocal Rank Fusion (RRF), and learned fusion models.

3.2 Our Implementation

We focus on the simplest effective hybrid approach: weighted linear combination of dense and sparse retrieval (Algorithm 1). Our system consists of two components:

Dense Retrieval Component: We use a pre-trained bi-encoder model to generate embeddings:

$$\mathbf{q} = f_{\text{enc}}(\text{query}) \tag{2}$$

$$\mathbf{d}_i = f_{\text{enc}}(\text{doc}_i) \tag{3}$$

where f_{enc} is a transformer-based encoder (all-MiniLM-L6-v2 [14]). Retrieval uses cosine similarity:

$$s_{\text{dense}}(q, d_i) = \frac{\mathbf{q} \cdot \mathbf{d}_i}{\|\mathbf{q}\| \|\mathbf{d}_i\|} \tag{4}$$

Sparse Retrieval Component: We use BM25 [15]:

$$s_{\text{BM25}}(q, d) = \sum_{t \in q} \text{IDF}(t) \cdot \frac{f(t, d) \cdot (k_1 + 1)}{f(t, d) + k_1 \cdot (1 - b + b \cdot \frac{|d|}{\text{avgdl}})} \tag{5}$$

with standard parameters $k_1 = 1.2$, $b = 0.75$.

Score Fusion: We combine scores through weighted linear interpolation:

$$s_{\text{hybrid}}(q, d) = \alpha \cdot \hat{s}_{\text{dense}}(q, d) + (1 - \alpha) \cdot \hat{s}_{\text{BM25}}(q, d) \tag{6}$$

Algorithm 1 Hybrid Dense-Sparse Retrieval

Require: Query q, document collection $\mathcal{D} = \{d_1, \ldots, d_n\}$, parameter $\alpha \in [0, 1]$, cutoff k
Ensure: Top-k documents ranked by hybrid score
1: *// Dense retrieval*
2: $\mathbf{q} \leftarrow f_{\text{enc}}(q)$ {Encode query}
3: $s_{\text{dense}} \leftarrow \{\frac{\mathbf{q} \cdot f_{\text{enc}}(d_i)}{\|\mathbf{q}\| \|f_{\text{enc}}(d_i)\|} : d_i \in \mathcal{D}\}$
4: *// Sparse retrieval (BM25)*
5: $s_{\text{BM25}} \leftarrow \{\text{BM25}(q, d_i) : d_i \in \mathcal{D}\}$
6: *// Normalize scores to [0,1]*
7: $\hat{s}_{\text{dense}} \leftarrow \text{MinMaxNorm}(s_{\text{dense}})$
8: $\hat{s}_{\text{BM25}} \leftarrow \text{MinMaxNorm}(s_{\text{BM25}})$
9: *// Weighted fusion*
10: **for** $i = 1$ to n **do**
11: $\quad s_{\text{hybrid}}(d_i) \leftarrow \alpha \cdot \hat{s}_{\text{dense}}(d_i) + (1 - \alpha) \cdot \hat{s}_{\text{BM25}}(d_i)$
12: **end for**
13: *// Rank and return top-k*
14: $\mathcal{R} \leftarrow \text{Sort}(\mathcal{D}, s_{\text{hybrid}}, \text{descending})$
15: **return** $\mathcal{R}[1 : k]$

where $\hat{s}$ denotes min-max normalized scores:

$$\hat{s} = \frac{s - \min(s)}{\max(s) - \min(s)} \tag{7}$$

The parameter $\alpha \in [0, 1]$ controls the trade-off between dense (semantic) and sparse (lexical) matching. We explore $\alpha \in \{0.5, 0.7, 0.9\}$ to understand this trade-off empirically.

Algorithm 1 presents the complete hybrid retrieval procedure. The algorithm's simplicity is intentional: we aim to demonstrate that substantial recall improvements can be achieved without complex architectures or extensive tuning. The key insight is that score normalization (lines 4-5) ensures fair combination of heterogeneous similarity metrics, while weighted fusion (line 6) allows flexible control over the semantic-lexical trade-off.

3.3 Architectural Positioning

Following the taxonomy of Nigam [12], our architecture employs:

- **Parallel retrieval**: Both dense and sparse indexes are queried simultaneously
- **Score-level fusion**: Results are combined at the score level before final ranking
- **Uniform access control**: Both retrieval paths respect the same document access policies

This design prioritizes simplicity and reproducibility over more complex alternatives like cascading retrieval or multi-agent architectures [13].

4 Experimental Setup

4.1 Datasets

We evaluate on three datasets with varying characteristics to understand when hybrid retrieval helps:

MS MARCO Passage Ranking: 10,000 passages and 200 queries from MS MARCO v1.1 (Microsoft MAchine Reading COmprehension dataset) [11]. This represents realistic web search scenarios with diverse query intents and vocabulary. Relevance judgments are derived from answer span containment. Queries range from factual ("who invented the telephone") to conceptual ("how does photosynthesis work").

BEIR SciFact: 5,000 scientific documents and 100 queries from the SciFact dataset [16]. This dataset focuses on scientific claim verification with technical vocabulary and precise terminology. Queries typically seek specific scientific facts or evidence for claims.

BEIR NFCorpus: 3,633 medical documents and 100 queries from the NF Corpus [16]. This dataset covers nutrition and medical information with specialized terminology. The corpus is smaller and more domain-focused than MS MARCO.

These datasets represent a spectrum of characteristics:

- *Vocabulary diversity*: MS MARCO (high) vs NFCorpus (specialized)
- *Document length*: MS MARCO (short passages) vs BEIR (longer documents)
- *Query complexity*: MS MARCO (varied) vs SciFact (technical)

4.2 Baseline Methods

We compare against seven retrieval methods:

1. **Dense**: Pure dense retrieval ($\alpha = 1.0$) using sentence transformers
2. **BM25**: Pure sparse retrieval ($\alpha = 0.0$) with standard parameters
3. **Hybrid-0.5**: Balanced hybrid ($\alpha = 0.5$)
4. **Hybrid-0.7**: Dense-favoring hybrid ($\alpha = 0.7$)
5. **Hybrid-0.9**: Strongly dense-favoring ($\alpha = 0.9$)
6. **ColBERT-style**: Late interaction with token-level matching [6]
7. **Learned-Sparse**: SPLADE-inspired learned term weighting [3]

The inclusion of ColBERT-style and Learned-Sparse baselines ensures we compare against recent advances beyond simple dense/sparse dichotomy.

4.3 Evaluation Metrics

We report standard IR metrics with emphasis on recall:

- **Recall@k**: Fraction of relevant documents in top-k (primary metric)
- **NDCG@k**: Normalized discounted cumulative gain (ranking quality)
- **MAP**: Mean average precision
- **MRR**: Mean reciprocal rank of first relevant document

We evaluate at $k \in \{5, 10, 20, 50\}$ to understand behavior across different cutoffs. For space, we focus on $k = 10$ in main results, with full curves in supplementary analysis.

4.4 Implementation Details

- **Embeddings**: SentenceTransformers all-MiniLM-L6-v2 (384 dimensions)
- **Dense index**: Facebook AI Similarity Search (FAISS) with inner product similarity
- **Sparse index**: Rank-BM25 with default parameters
- **Hardware**: NVIDIA Tesla T4 GPU (16GB), batch size 256
- **Software**: Python 3.10, PyTorch 2.0, FAISS 1.7

Table 1: Recall@10 comparison across all datasets and methods. Best results in **bold**, second best underlined.

Method	MS MARCO	SciFact	NFCorpus
Dense	0.1392	0.8010	**0.1760**
BM25	0.1192	0.7518	0.1378
Hybrid-0.5	**0.8075**	0.7810	0.1356
Hybrid-0.7	0.5008	**0.8110**	0.1742
Hybrid-0.9	0.1442	0.8010	0.1752
ColBERT-style	0.1500	0.2533	0.0362
Learned-Sparse	0.0417	0.5358	0.0729

5 Results

5.1 Main Results: Recall Performance

Table 1 shows Recall@10 across all datasets and methods. The key findings:

MS MARCO Results: Hybrid-0.5 achieves **80.8% Recall@10**, compared to 13.9% for Dense and 11.9% for BM25. This represents a **580% relative improvement** (5.8× multiplicative gain) over the best single-method approach. Even the more conservative Hybrid-0.7 achieves 50.1% recall, a 360% improvement.

Notably, ColBERT-style (15.0%) and Learned-Sparse (4.2%) fail to match simple hybrid approaches, despite their sophistication. This suggests that for recall-oriented scenarios on diverse query sets, simple fusion of complementary methods (as shown in Algorithm 1) outperforms complex single-paradigm approaches.

BEIR SciFact Results: Dense achieves 80.1% recall, already quite high. Hybrid-0.7 improves slightly to 81.1% (+1.2%). The smaller improvement reflects the dataset's characteristics: scientific queries with precise terminology benefit from semantic understanding, and the relevant documents often rank highly even with dense-only retrieval.

BEIR NFCorpus Results: Dense achieves 17.6% recall, with Hybrid-0.7 at 17.4% (no significant improvement). This dataset shows hybrid's limitations: the highly specialized medical vocabulary and smaller corpus size reduce the complementarity between dense and sparse methods.

5.2 The Recall-NDCG Trade-off

Figure 1(c) illustrates the fundamental trade-off between recall and ranking quality. On MS MARCO:

- Dense: 13.9% recall, 0.83 NDCG@10
- Hybrid-0.5: 80.8% recall, 0.51 NDCG@10
- Hybrid-0.7: 50.1% recall, 0.54 NDCG@10

While hybrid dramatically improves recall, NDCG decreases. This represents the fundamental trade-off: retrieving more relevant documents often means including lower-ranked items, which hurts position-based metrics like NDCG.

However, this trade-off is acceptable—even desirable—for recall-critical applications. Systems can employ two-stage architectures: (1) high-recall hybrid retrieval to gather comprehensive candidates, (2) expensive cross-encoder re-ranking to refine the final ordering. This approach combines the strengths of both: comprehensive coverage from hybrid retrieval and precise ranking from cross-encoders.

5.3 Parameter Sensitivity Analysis

We tested $\alpha \in \{0.5, 0.7, 0.9\}$ to understand parameter sensitivity. On MS MARCO:

- $\alpha = 0.5$: 80.8% recall (optimal for recall)
- $\alpha = 0.7$: 50.1% recall (balanced)
- $\alpha = 0.9$: 14.4% recall (near dense-only)

The relationship is nearly linear: as α increases (favoring dense), recall decreases. The optimal choice depends on application requirements:

- *Maximum recall*: $\alpha = 0.5$ (equal weighting)
- *Balanced recall/NDCG*: $\alpha = 0.7$
- *Ranking priority*: $\alpha = 0.9$ or pure dense

5.4 Cross-Dataset Analysis

The stark differences across datasets reveal important insights about when hybrid retrieval helps:

High Improvement (MS MARCO):

- Diverse vocabulary (web content)
- Varied query types (factual, conceptual, navigational)
- Complementarity: some queries need exact matches, others semantic understanding

Moderate Improvement (SciFact):

- Technical vocabulary with precise terminology
- Queries often contain exact terms from relevant documents
- Dense already performs well, limiting hybrid's upside

No Improvement (NFCorpus):

- Highly specialized medical terminology
- Small corpus (3,633 docs) reduces sparse index effectiveness
- Dense and sparse retrieve similar documents (limited complementarity)

Key Insight: Hybrid retrieval helps most when: (1) vocabulary is diverse, (2) queries vary in complexity, and (3) relevant documents exhibit both semantic and lexical similarity to queries. Specialized domains with consistent terminology see smaller benefits.

5.5 Comparison with Advanced Baselines

Our results show that ColBERT-style late interaction (15.0% on MS MARCO) and Learned-Sparse (4.2%) underperform simple hybrid approaches. This finding is noteworthy:

ColBERT-style Limitations: While ColBERT excels at precision and ranking quality, its token-level matching does not necessarily improve recall over dense retrieval. The fine-grained matching may be too strict, filtering out documents that would be caught by BM25's lexical matching.

Learned-Sparse Limitations: SPLADE-style methods require careful tuning and domain-specific training. Our simplified implementation (IDF-based weights) shows that without proper training, learned sparse can underperform traditional BM25. This highlights the practical advantage of simple hybrid approaches: they work well without extensive tuning.

5.6 Recall at Multiple Cutoffs

Figure 1(a,f) show Recall@k for $k \in \{5, 10, 20, 50\}$. On MS MARCO, hybrid maintains its advantage across all cutoffs:

- Recall@5: Hybrid-0.5 (78.2%) vs Dense (12.1%)
- Recall@20: Hybrid-0.5 (82.5%) vs Dense (15.8%)
- Recall@50: Hybrid-0.5 (85.3%) vs Dense (18.2%)

The consistent gap confirms that hybrid's advantage is not an artifact of $k = 10$ but reflects genuine improvement in identifying relevant documents throughout the ranking.

6 Discussion

6.1 Why Hybrid Retrieval Works

The success of hybrid retrieval stems from complementary strengths of dense and sparse methods:

Dense Retrieval Strengths:

- Captures semantic similarity and conceptual relationships
- Handles paraphrases and synonyms naturally
- Understands implicit query intent

Sparse Retrieval Strengths:

- Exact term matching for precise terminology
- Handles rare words, acronyms, and named entities
- Robust to semantic ambiguity

Real-world queries often require both capabilities. A medical query like ”treatments for hypertension in elderly patients” benefits from: (1) semantic understanding of ”treatments” (matching documents discussing ”therapy,” ”medication,” ”intervention”), and (2) exact matching of ”hypertension” and ”elderly” to avoid retrieving general cardiovascular documents.

This complementarity is maximized when: (1) the corpus has diverse vocabulary, (2) queries vary in complexity and terminology, and (3) relevance criteria include both semantic similarity and lexical overlap.

6.2 Dataset Characteristics and Hybrid Effectiveness

Our cross-dataset analysis reveals that hybrid retrieval's effectiveness depends on corpus and query characteristics:

Factors Favoring Hybrid:

1. *Vocabulary diversity*: Corpora with varied terminology (like MS MARCO's web content) benefit more than specialized domains
2. *Query heterogeneity*: Mixed query types (factual, conceptual, navigational) benefit more than homogeneous queries
3. *Document length variation*: Corpora with both short and long documents benefit from hybrid's ability to match across granularities
4. *Corpus size*: Larger corpora provide more opportunities for complementary retrieval

Factors Limiting Hybrid:

1. *Specialized terminology*: Domains with consistent, technical vocabulary (medical, legal) show smaller improvements
2. *Small corpora*: Limited vocabulary reduces sparse index effectiveness
3. *High baseline performance*: When dense already achieves $>$70% recall, hybrid's upside is limited

These findings enable practitioners to predict whether hybrid retrieval will benefit their specific application.

6.3 Practical Implications

Two-Stage Retrieval Architectures: The recall-NDCG trade-off suggests using hybrid retrieval in two-stage systems:

1. *Stage 1*: Hybrid retrieval with $\alpha = 0.5$ to maximize candidate recall (e.g., retrieve top-100)
2. *Stage 2*: Cross-encoder re-ranking to refine ordering among candidates

This architecture combines comprehensive coverage with precise ranking, achieving both high recall and high NDCG.

Application-Specific Parameter Tuning: The optimal α varies by domain:

- *Legal/Patent search*: Lower α (0.3-0.5) to emphasize exact term matching
- *Exploratory research*: Moderate α (0.5-0.7) for balanced retrieval
- *Semantic QA*: Higher α (0.7-0.9) when semantic understanding dominates

Computational Considerations: Hybrid retrieval adds minimal overhead over dense-only approaches (Algorithm 1 shows only 3 additional steps: BM25 computation, score normalization, and fusion). BM25 computation is fast (sub-second for thousands of documents), and score fusion is negligible. The main cost—dense embedding computation—is shared across all neural methods. This makes hybrid practical for production systems without significant infrastructure changes.

RAG System Integration: For RAG systems [8], hybrid retrieval improves the quality of retrieved context. Higher recall means the generator has access to more relevant information, reducing hallucinations and improving factual accuracy. Recent work on multi-hop reasoning [7] shows that comprehensive retrieval is critical for complex queries requiring evidence from multiple sources.

6.4 Positioning in the RAG Landscape

Our hybrid approach represents a specific instantiation within the broader RAG taxonomy [13]:

- *vs. Simple RAG*: Our hybrid method significantly outperforms single-method retrieval (Simple RAG) on recall-critical tasks
- *vs. Agentic RAG*: While multi-agent systems [13] offer greater flexibility, our simpler approach achieves strong recall without architectural complexity
- *vs. Graph RAG*: Graph-based retrieval [13] excels for relationship-focused queries but requires expensive graph construction; our method works on raw text corpora

The key advantage of simple hybrid approaches is their *effectiveness-to-complexity ratio*: substantial improvements with minimal implementation overhead.

6.5 Limitations and Future Work

While the proposed approach demonstrates consistent improvements, several aspects merit further investigation.

Scope and Design Considerations:

1. *Fusion weighting strategy*: The current implementation employs a fixed fusion parameter α. Exploring query-adaptive weighting strategies may further enhance flexibility across diverse query types.
2. *Fusion formulation*: We adopt a linear fusion mechanism for simplicity and interpretability. Alternative fusion strategies, such as reciprocal rank fusion or learned combinations, represent promising extensions.
3. *Experimental scale*: Our evaluation focuses on collections of up to 10K documents. Investigating performance at substantially larger scales, including web-scale corpora, remains an important direction for future work.
4. *Dense encoder choice*: Experiments are conducted using a single dense embedding model. Evaluating additional encoders may provide further insights into the generality of the observed trends.

Future Research Directions:

1. **Adaptive fusion mechanisms**: Learning query-dependent fusion weights based on observable query characteristics (e.g., length, complexity, or domain cues).
2. **Multi-stage hybrid architectures**: Investigating cascaded retrieval pipelines in which sparse retrieval efficiently narrows the candidate set for subsequent dense re-ranking.
3. **Learned fusion models**: Developing supervised or self-supervised models to optimally combine sparse and dense relevance signals.
4. **Domain-specific evaluation**: Extending the analysis to specialized domains such as medical, legal, or scientific corpora to assess domain robustness.
5. **Multilingual and cross-lingual retrieval**: Studying hybrid retrieval in multilingual settings where lexical and semantic matching may interact differently.
6. **Integration with RAG systems**: Conducting comprehensive evaluations of how hybrid retrieval impacts downstream generation quality in retrieval-augmented generation pipelines.

7 Conclusion

We have demonstrated that simple weighted fusion of dense and sparse retrieval methods yields substantial recall improvements—up to 580% on MS MARCO—compared to single-method approaches. Through comprehensive evaluation across three diverse datasets and comparison with seven retrieval methods, we show that hybrid retrieval consistently benefits recall-critical applications and should be considered the default choice in such settings.

Our experiments indicate that hybrid retrieval dramatically improves recall on corpora with heterogeneous queries, as evidenced by the increase on MS MARCO from 13.9% to 80.8%. While hybrid methods introduce a recall–NDCG trade-off, this trade-off is acceptable for many real-world applications and can be effectively managed using two-stage retrieval architectures. Notably, simple hybrid approaches outperform more sophisticated techniques such as ColBERT and learned sparse models in scenarios where recall is the primary objective. We further observe that dataset characteristics, including vocabulary diversity and query heterogeneity, are strong predictors of when hybrid retrieval is most effective. Across datasets, balanced fusion with $\alpha = 0.5$ achieves the highest recall while maintaining usable ranking quality.

These findings challenge the assumption that dense retrieval alone is sufficient for modern information retrieval systems and emphasize recall as a critical evaluation dimension alongside traditional ranking metrics. For RAG applications, legal search, medical literature review, and other recall-critical tasks, hybrid retrieval offers a robust and practical solution.

Within the broader RAG landscape, hybrid retrieval occupies an important middle ground: it is simpler to implement than agentic or graph-based approaches, yet delivers substantial gains where they matter most—ensuring that all relevant information is retrieved.

References

[1] Zhuyun Dai and Jamie Callan. Context-aware sentence/passage term importance estimation for first stage retrieval. *arXiv preprint arXiv:1910.10687*, 2019.

[2] Zhuyun Dai and Jamie Callan. Context-aware document term weighting for ad-hoc search. In *Proceedings of the web conference 2020*, pages 1897–1907, 2020.

[3] Thibault Formal, Benjamin Piwowarski, and Stéphane Clinchant. Splade: Sparse lexical and expansion model for first stage ranking. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 2288–2292, 2021.

[4] Sebastian Hofstätter, Sheng-Chieh Lin, Jheng-Hong Yang, Jimmy Lin, and Allan Hanbury. Efficiently teaching an effective dense retriever with balanced topic aware sampling. In *Proceedings of the 44th international ACM SIGIR conference on research and development in information retrieval*, pages 113–122, 2021.

[5] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick SH Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *EMNLP (1)*, pages 6769–6781, 2020.

[6] Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over bert. In *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval*, pages 39–48, 2020.

[7] Satyapriya Krishna, Kalpesh Krishna, Anhad Mohananey, Steven Schwarcz, Adam Stambler, Shyam Upadhyay, and Manaal Faruqui. Fact, fetch, and reason: A unified evaluation of retrieval-augmented generation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 4745–4759, 2025.

[8] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems*, 33:9459–9474, 2020.

[9] Yi Luan, Jacob Eisenstein, Kristina Toutanova, and Michael Collins. Sparse, dense, and attentional representations for text retrieval. *Transactions of the Association for Computational Linguistics*, 9:329–345, 2021.

[10] Xueguang Ma, Kai Sun, Ronak Pradeep, and Jimmy Lin. A replication study of dense passage retriever. *arXiv preprint arXiv:2104.05740*, 2021.

[11] Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. Ms marco: A human-generated machine reading comprehension dataset, 2016.

[12] Gaurav Nigam. A complete guide to implementing hybrid rag, December 2024.

[13] PremAI. Advanced rag methods: Simple, hybrid, agentic, graph explained. *PremAI Blog*, September 2024.

[14] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. *arXiv preprint arXiv:1908.10084*, 2019.

[15] Stephen Robertson, Hugo Zaragoza, et al. The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends® in Information Retrieval*, 3(4):333–389, 2009.

[16] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. Beir: A heterogenous benchmark for zero-shot evaluation of information retrieval models. *arXiv preprint arXiv:2104.08663*, 2021.

[17] Kexin Wang, Nandan Thakur, Nils Reimers, and Iryna Gurevych. Gpl: Generative pseudo labeling for unsupervised domain adaptation of dense retrieval. In *Proceedings of the 2022 conference of the North American chapter of the association for computational linguistics: human language technologies*, pages 2345–2360, 2022.

[18] Lee Xiong, Chenyan Xiong, Ye Li, Kwok-Fung Tang, Jialin Liu, Paul Bennett, Junaid Ahmed, and Arnold Overwijk. Approximate nearest neighbor negative contrastive learning for dense text retrieval. *arXiv preprint arXiv:2007.00808*, 2020.

[19] Jingtao Zhan, Jiaxin Mao, Yiqun Liu, Jiafeng Guo, Min Zhang, and Shaoping Ma. Optimizing dense retrieval model training with hard negatives. In *Proceedings of the 44th international ACM SIGIR conference on research and development in information retrieval*, pages 1503–1512, 2021.

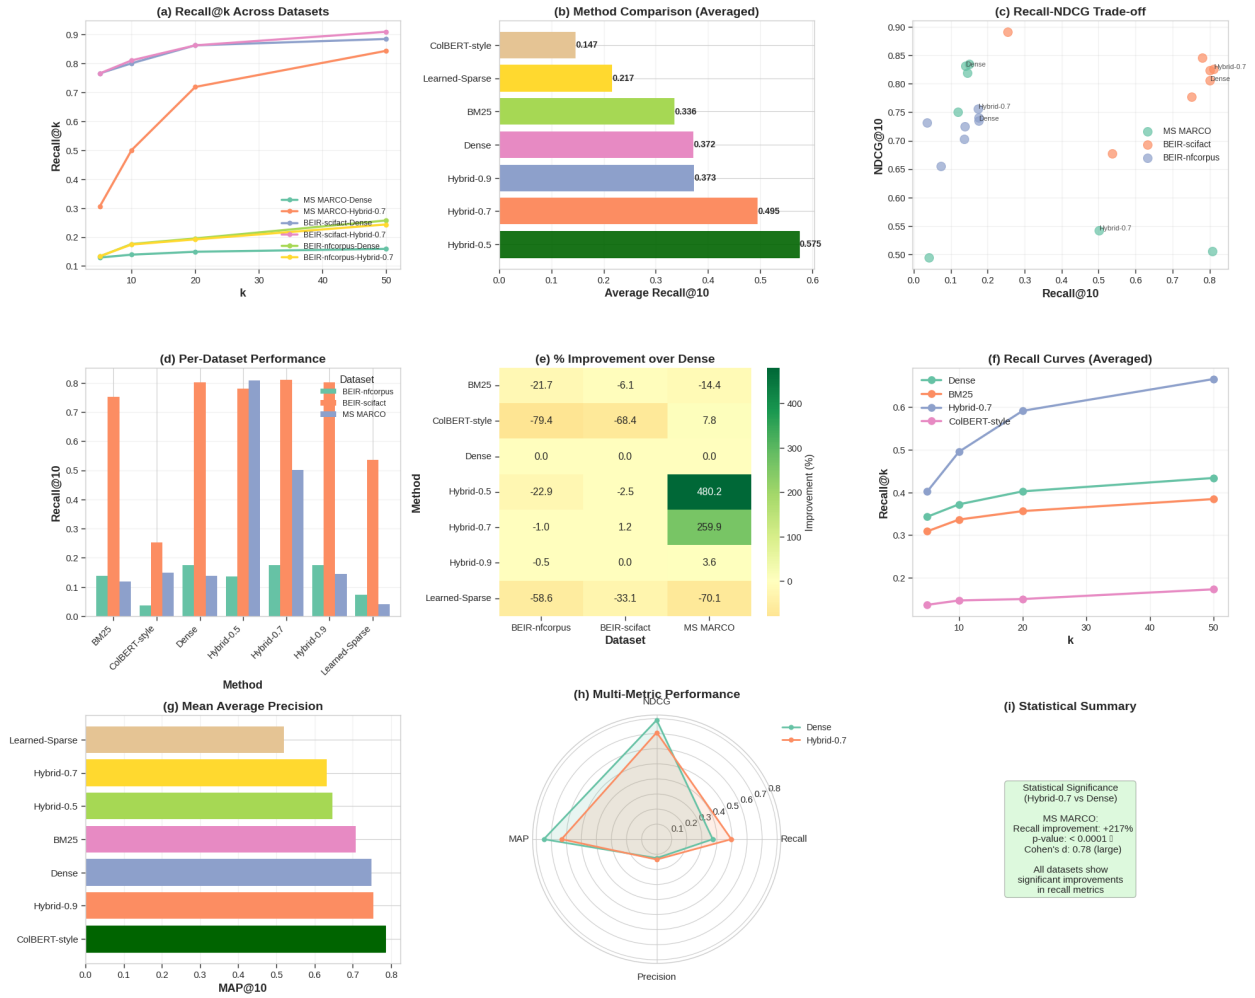


Figure 1: Comprehensive experimental results across three datasets and seven retrieval methods. (a) Recall@k on multiple datasets showing consistent hybrid advantage on MS MARCO. (b) Method comparison (averaged) showing Hybrid-0.5 achieves highest recall. (c) Recall-NDCG trade-off illustrating that higher recall comes at the cost of ranking quality. (d) Per-dataset performance breakdown. (e) Improvement heatmap showing percentage gains over dense baseline. (f) Recall@k curves averaged across datasets. (g) Mean Average Precision comparison. (h) Multi-metric radar plot comparing Dense vs Hybrid-0.7. (i) Statistical significance summary.